

HOMWORK EXERCISE-2 (FOR LECTURE 7)

This homework covers Lectures 7-10. Please take a look at the material for the context and notation. You have three weeks to complete the homework.

The homework requires you to study minimax problems in the context of Generative Adversarial Networks (GANs). Since the state-of-the-art GAN training is computationally expensive, you will implement a simple GAN with linear generator and a dual variable. Then, you will implement two variants of the stochastic gradient descent, AMSGrad and RMSProp. You will use these algorithms to train a neural network to solve the image classification task. For concrete demonstrations, we will use one of the most popular deep learning frameworks, called “PyTorch”. These problems should expose you the fundamentals of GAN training at a basic level.

For the homework, we provide some basic code upon which you need to build functionality. You should only add code in the sections marked with the `TODO` comment, but feel free to modify other parts of the code if you really need to. Please read the `README.md` file for further required python modules.

1 Minimax problems and GANs

1.1 Theoretical recap: stationary points and convergence in minmax games

Consider a stylized function $f : \mathbb{R}^2 \rightarrow \mathbb{R}$, such that we have $f(x, y) = (x - 1)(y - 1)$.

- (10 points) Write down the first-order stationary points of f , and classify them as local minimum, local maximum, or saddle point by inspecting its Hessian.
- (10 points) Show that $(x^*, y^*) = (1, 1)$ is a solution to the minimax problem $\min_x \max_y f(x, y)$. You can quantify the solution by using the following saddle point inequality: $f(x^*, y^*) \geq f(x^*, y)$ and $f(x^*, y^*) \leq f(x, y^*)$, for all x, y .
- (20 points) One possible attempt at finding this solution via iterative first-order methods is to perform gradient updates on the variables x and y . More precisely for $\gamma > 0$ consider the simultaneous gradient descent/ascent updates

$$x_{k+1} = x_k - \gamma \nabla_x f(x_k, y_k), \quad y_{k+1} = y_k + \gamma \nabla_y f(x_k, y_k)$$

Show that the sequence of iterates $\{x_k, y_k\}_{k=0}^\infty$ starting from any point $(x_0, y_0) \neq (1, 1)$ diverges, for any $\gamma > 0$. Find the rate at which the distance of the sequence $\{x_k, y_k\}$ to the origin grows as the number of iterations k increases.

1.2 Practice: GAN

Please review Lecture 10 for the notation and the setup in this exercise. As a basic example for the implementation of GANs, we will try to learn the density of a mixture of Gaussian (MoG) distribution from empirical samples using the Wasserstein GAN (WGAN) framework. MoG distribution is multi-modal, and hence, we can visualize multiple “modes” of the distribution.

Both your generator neural network h_x and dual variable neural network d_y classes are defined as single hidden layer networks:

$$\mathcal{H} := \{h : h_x(\omega) = X_2 \text{relu}(X_1 \omega + x_1) + x_2\}, \quad \mathcal{D} := \{d : d_y(\mathbf{a}) = Y_2 \text{relu}(Y_1 \mathbf{a} + y_1) + y_2\}, \quad (1)$$

where $\mathbf{x} = [x_1; x_2; X_1; X_2]$ are the “generator” parameters, and $\mathbf{y} = [y_1; y_2; Y_1; Y_2]$ are the “dual” or the “discriminator” parameters. The dimensions of these parameters will be apparent from the context as well as the base code provided along with the homework.

- (10 points) Fill in the missing code in the file `variables.py` implementing the two neural networks and the spectral normalization (Myato 2018) method that enforces a Lipschitz constraint. For this you have to do the following:
 - Create the weights and any auxilliary variables you might need for the Spectral Normalization method (**Hint:** Think carefully which variables need to be updated via autograd and hence need to be implemented as `torch.nn.Modules` or `torch.nn.Parameters` and which variables you will need to use `.detach()` and manual updates)

(b) Fill in the `enforce_lipschitz` function which performs Spectral Normalization as presented in Miyato 2018

Then, implement a stochastic estimate of the objective function of the minimax game (in `trainer.py`):

$$\min_{\mathbf{x} \in \mathcal{X}} \max_{\mathbf{y} \in \mathcal{Y}} E[d_{\mathbf{y}}(\mathbf{a})] - E[d_{\mathbf{y}}(h_{\mathbf{x}}(\omega))] = \min_{h \in \mathcal{H}} \max_{d \in \mathcal{F}} E[d(\mathbf{a})] - E[d(h(\omega))], \quad (2)$$

where $\mathbf{a} \sim \hat{\mu}_n$ is distributed with respect to the empirical estimate $\hat{\mu}_n$ of the true distribution μ^{\dagger} (i.e., it is distributed uniformly on a set of samples drawn from the true distribution), and $\omega \sim p_{\Omega}$ is the noise distribution that we push-forward with the generator neural network $h_{\mathbf{x}}(\omega) \sim h_{\mathbf{x}} \# p_{\Omega}$. The stochastic estimate that you should implement simply corresponds to an average over finite samples drawn from the target and noise distribution. Use the methods available for the class `torch.distributions.Distribution` from the PyTorch module.

Remark: The `enforce_lipschitz` does not need to return a value. It should modify the parameters \mathbf{y} of the dual network $d_{\mathbf{y}}$ in situ. In order to do this we recommend directly modifying any tensor's values by rewriting their data attribute e.g., `y.data = y.data + 1`. This avoids PyTorch's automatic tracking of operations for automatic differentiation, which *might* cause issues.

- (10 points) Compare and contrast Spectral Normalization with Spectral *Regularization* (mentioned in the paper) and also compare and contrast with other methods of Lipschitz constraints used in WGAN training, namely Weight Clipping and Gradient penalty. Write an explanation of the differences.
- (25 points) Complete the missing functions in `trainer.py`. You should implement an alternating gradient ascent/descent update, training the generator 1 time for every 5 dual updates. More specifically, you will implement the conceptual algorithm below using the parameters of the neural networks

$$\mathbf{d}^{k+1} = \text{clip}(\mathbf{d}^k + \gamma \text{SG}_d(\mathbf{d}^k, h^k k)), \quad (\text{always}) \quad (3)$$

$$h^{k+1} = h^k - \gamma \text{SG}_h(\mathbf{d}^{k+1}, h^k) \quad (\text{if } k \bmod 5 = 0), \quad (4)$$

where SG is the stochastic gradient oracle. To perform the optimization, use 2 separate instances of Adam with $(\beta_0, \beta_1) = (0.0, 0.9)$ instead of the simpler SGD steps used for illustrative purposes above.

After filling in your code, run the script `train.py`. Include the generated plots in your report and comment on your findings.

2 Optimizers of Neural Network

The goal of this exercise is to implement different optimizers for a handwritten digit classifier using the well-known MNIST dataset. This dataset has 60000 training images, and 10000 test images. Each image is of size 28×28 pixels, and shows a digit from 0 to 9.

- General guideline: Complete the codes marked with *TODO* in the `optimizers.py`.
- The implementation of mini-batch SGD and SGD with HB momentum are provided as examples.

Vanilla Minibatch SGD
Input: learning rate γ
1. initialize θ_0
2. For $t = 0, 1, \dots, N-1$:
obtain the minibatch gradient $\hat{\mathbf{g}}_t$
update $\theta_{t+1} \leftarrow \theta_t - \gamma \hat{\mathbf{g}}_t$

Minibatch SGD with Momentum
Input: learning rate γ , momentum ρ
1. initialize $\theta_0, \mathbf{m}_0 \leftarrow \mathbf{0}$
2. For $t = 0, 1, \dots, N-1$:
obtain the minibatch gradient $\hat{\mathbf{g}}_t$
update $\mathbf{m}_{t+1} \leftarrow \rho \mathbf{m}_t + \hat{\mathbf{g}}_t$
update $\theta_{t+1} \leftarrow \theta_t - \gamma \mathbf{m}_{t+1}$

- Implement of following optimizers: